Applied Forecasting Algorithms

Probabilistic Flight Delay Forecasting

with Cross-Airport Delay Propagation Analysis

Dhirubhai Ambani University Technology

Formerly DA-IICT

Group Members:

Aarsh Adhvaryu · Nikita Sharma · Ram Sharma

202518031 · 202518038 · 202518009

Course Name	Applied Forecasting Algorithms ·
Instructor Name	Prof. Pritam Anand
Institution Name	Dhirubhai Ambani University
Submission Date	5 May 2026
Live App	https://flightdelayforecasting.streamlit.app
GitHub Repo	github.com/nikitasharma120/flight_delay_forecasting

1. Abstract

Flight delays at US hub airports are not isolated events, they cascade across the network via shared aircraft rotations, crew, and passengers. Current operational forecasting systems issue single point estimates and treat airports independently, providing no uncertainty quantification and ignoring the propagation structure that makes hub delays correlated in time and space.

This paper presents a reproducible, open-data probabilistic forecasting pipeline applied to **22,085,189** raw US BTS On-Time Performance records (January 2023 – May 2025) across the 10 busiest US hub airports. The pipeline proceeds in five stages: (1) baseline benchmarking; (2) statistical modelling (ARIMA, SARIMA, SARIMAX); (3) deep learning point forecasting (LSTM, GRU, Tube Loss LSTM); (4) probabilistic calibration via Quantile Regression Neural Networks (QRNN) with Pinball Loss; and (5) cross-airport propagation analysis combining tail-number flight chains with Granger causality testing.

Key findings: residuals are non-Gaussian at every airport under every model (Shapiro-Wilk $p < 10^{-28}$), invalidating standard Gaussian prediction intervals. SARIMA(1,0,2)(0,1,1)[7] achieves the best point forecast RMSE of **4.896 min** at PHX — a 23.5% improvement over the best baseline. QRNN achieves prediction interval coverage probability (PICP) of 0.717–0.855 across airports. Cross-airport Granger causality identifies **68 of 90** hub pairs as significant ($p < 0.05$), and tail-number chain analysis quantifies a mean delay carry-over fraction of **74.2%** across 793,941 consecutive flight legs. A novel hourly alert system achieves macro-average F1 = **0.760** on a 366-day held-out test — a 2.1× improvement over a daily-granularity system.

Problem Statement: **Build a probabilistic forecasting pipeline providing calibrated prediction intervals and modelling cross-airport delay propagation for early disruption warnings.**

Methods Used: **ARIMA, SARIMA, SARIMAX, LSTM, GRU, Tube Loss LSTM, QRNN (Pinball Loss), Granger causality, tail-number chain analysis, Streamlit dashboard.**

Key Findings: **(1) Universal non-Gaussianity motivates QRNN. (2) SARIMA champion model. (3) 68/90 propagation pairs significant; hourly alert F1 = 0.760.**

2. Introduction

2.1 Background and Motivation

Context of the Forecasting Problem: The US domestic aviation system handles approximately 45,000 flights daily across a network of interconnected hubs. Departure delays at one hub propagate physically through the network as the same aircraft, crew, and connecting passengers travel downstream. This propagation mechanism — commonly called the "rotation effect" — means that a 45-minute delay at Chicago O'Hare at 08:00 can produce a 35-minute delay at Dallas–Fort Worth by 13:00, and further cascades to Miami by 17:00. The US Bureau of Transportation Statistics (BTS) independently labels this mechanism as *LATE_AIRCRAFT_DELAY* in its per-flight records, providing a ground-truth signal for validation.

Why It Is Important and Its Impact: The FAA estimates that airline delays cost the US economy approximately $33 billion annually. Passenger disruption, crew misplacement, gate conflicts, and fleet downgrading all compound primary delay costs. Crucially, these costs are nonlinear: a 5-minute departure delay typically causes 0–2 minutes of network impact, but a 60-minute delay at a major hub can cascade into hundreds of downstream minutes across dozens of flights. Accurate and calibrated probabilistic forecasts that quantify uncertainty enable operations centres to pre-position standby crew, issue gate-hold advisories, and notify connecting passengers before cascades materialise.

What Are the Associated Challenges: Three challenges motivate this project. First, **non-Gaussian noise**: flight delay distributions are zero-bounded, heavily right-skewed, and multi-modal, rendering standard Gaussian prediction intervals systematically over-narrow at the right tail where operational impact is greatest. Second, **network dependence**: standard time-series models treat each airport independently, discarding the causal information contained in upstream delays. Third, **granularity mismatch**: daily-average models miss intra-day patterns (morning weather, afternoon thunderstorm, evening cascade) that are the actual operational signal.

2.2 Problem Statement

We address the following composite forecasting problem: given the historical daily mean departure delay time series at 10 US hub airports, together with the flight-level aircraft rotation graph, produce (a) calibrated 90% prediction intervals for the next-day average departure delay at each hub and (b) an hourly alert system that predicts, with a precision exceeding 0.80, which downstream hub airports will experience delay spikes within 1–7 days of an observed upstream elevation.

2.3 Objectives

- Benchmark four naive baselines and establish a minimum acceptable RMSE floor.
- Fit and evaluate ARIMA, SARIMA, and SARIMAX on the daily delay series per airport, with rigorous residual diagnostics.
- Train LSTM, GRU, and Tube Loss LSTM deep learning models and compare against statistical baselines.
- Train a QRNN with Pinball Loss to produce calibrated 90% prediction intervals; evaluate with PICP, PINAW, and Winkler score.
- Formally prove non-Gaussianity of residuals across all model families and airports.

- Quantify physical delay propagation via tail-number flight chains and validate against BTS LATE_AIRCRAFT_DELAY.
- Identify statistically significant propagation relationships via Granger causality on daily aggregated series.
- Build and evaluate an hourly alert system leveraging Granger lags converted to hourly windows.
- Deploy all findings in an interactive Streamlit operations dashboard with eight functional pages.

3. Data Source

Source: **US Bureau of Transportation Statistics (BTS) On-Time Performance Dataset**

URL: *transtats.bts.gov — free, no API key, no trial, available since 1987*

Time Period: **January 2023 – May 2025 (29 monthly CSV/Parquet downloads)**

Raw Volume: **22,085,189 individual flight records across 39 fields**

Each record represents one scheduled departure and includes: FlightDate, Reporting_Airline, Tail_Number (aircraft registration), Origin, Dest, CRS_DEP_TIME (scheduled), DEP_TIME (actual), DEP_DELAY (minutes), ARR_DELAY, CANCELLED, DIVERTED, and the five BTS delay-cause breakdown fields: CARRIER_DELAY, WEATHER_DELAY, NAS_DELAY, SECURITY_DELAY, and LATE_AIRCRAFT_DELAY. The LATE_AIRCRAFT_DELAY field is particularly valuable because BTS independently labels how many minutes of each departure delay were caused by a late inbound aircraft, providing a ground-truth cross-validation signal for our tail-number propagation analysis.

3.1 Data Characteristics

- Variables: 39 columns per flight record (continuous delay times, categorical carrier/origin/dest codes, binary flags)
- Granularity: individual flight level (minute precision), aggregated to daily mean per airport for modelling
- Frequency: daily aggregation yields 8,819 rows × 10 airports after hub filtering
- Target variable: daily mean departure delay in minutes per airport (avg_dep_delay)
- BTS defines a delay as DEP_DELAY ≥ 15 minutes; pct_delayed = fraction of flights meeting this threshold
- Airport scope: 10 busiest hubs by departure volume — ATL (1,036K), DEN (951K), DFW (941K), ORD (894K), CLT (631K), LAX (605K), PHX (591K), LAS (590K), SEA (516K), MCO (501K)

3.2 Data Pre-processing

Missing Values Handling

- Drop cancelled (CANCELLED=1) and diverted (DIVERTED=1) flights: −372,559 rows (1.69%)
- Drop records with null DEP_DELAY: no additional rows lost post-cancellation filter
- Fill all delay-cause NaN fields (CARRIER_DELAY, WEATHER_DELAY, NAS_DELAY, SECURITY_DELAY, LATE_AIRCRAFT_DELAY) with 0 — NaN in these fields means the delay cause was not applicable, not missing
- Clip DEP_DELAY to [−60, +600] minutes: removes 32,271 data-entry artefacts (e.g., impossible −90 min “delay” or +900 min outliers)

Normalisation / Scaling

- No global normalisation is applied to the target series for ARIMA/SARIMA (models operate in original units: minutes)
- Deep learning inputs: features are standardised per airport using StandardScaler fit on training data only; inverse transform is applied before evaluation to report metrics in original minutes
- QRNN quantile outputs are inverse-transformed to minutes before PICP/PINAW/Winkler computation

Feature Engineering

- 23 features total per airport per day:
- **Cyclical encodings**: sin/cos pairs for month-of-year, day-of-week, and day-of-year. These preserve the cyclic adjacency (December is adjacent to January) that one-hot encoding discards.
- **Lag features**: lag_1 (yesterday), lag_2, lag_3, lag_7 (same weekday last week), lag_14. All use .shift(1) to guarantee no future leakage.
- **Rolling statistics**: rolling mean at 7, 14, 30 days; rolling std at 7 and 14 days. Computed with .shift(1).rolling(window) to prevent look-ahead bias.
- **Calendar flags**: is_weekend, is_monday, is_friday (high-delay days per EDA), is_summer (June–August: thunderstorm season), is_holiday_season (November–December: Thanksgiving and Christmas peaks).
- **Operational features**: n_flights (total departures that day), pct_delayed (fraction of flights ≥15 min delayed).

Exploratory Analysis

Monthly × airport delay heatmaps reveal three systematic patterns: a weekly cycle (Friday, Sunday, Monday highest), a summer thunderstorm peak in June–August (MCO July = 31.5 min, DFW July = 24.3 min), and a November–December holiday spike. Delay cause decomposition shows Late Aircraft Delay (orange) is the dominant cause at most airports, constituting over 40% of total average delay at CLT, ORD, and MCO. This directly motivates the propagation analysis in Section 6.

STL decomposition across all 10 airports reveals seasonal strength $F_s \leq 0.186$ — all below the 0.64 threshold for meaningful seasonality. Trend dominates (F_t = 0.215–0.500). Despite weak F_s, ACF analysis confirms statistically significant weekly lags at 7, 14, 21, and 28 days, justifying the seasonal period m=7 in SARIMA. The high residual standard deviation (4.0–14.7 min) confirms that irregular shocks (storms, incidents, the 2024-07-19 CrowdStrike outage) dominate variance, establishing the need for probabilistic uncertainty quantification.

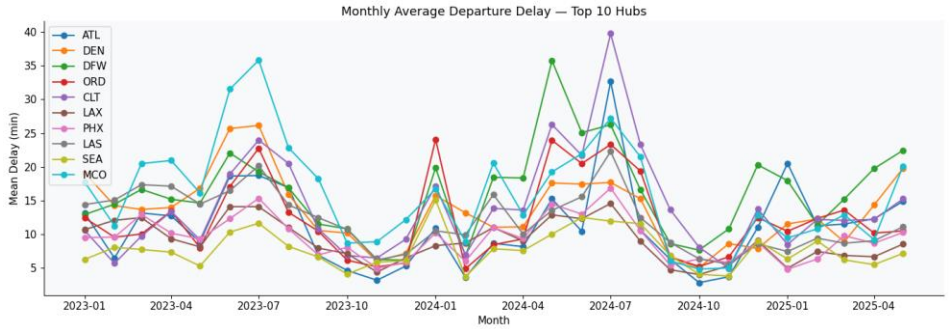

4. Methodology

Forecasting Horizon: **1 day ahead (daily mean departure delay) — horizon = 1**

Input–Output Structure: **For deep learning: input = sequence of last SEQ_LEN=21 days × 23 features; output = scalar daily mean delay. For statistical: walk-forward one-step-ahead evaluation over 366 test days.**

Train–Test Split: **Chronological, no shuffling. Train = 2023-01-01 to 2024-05-31 (516 days); Test = 2024-06-01 to 2025-05-31 (366 days). Assertion added (FIX4) to confirm all airports share identical last date.**

Reproducibility: **SEED=42 fixed in Python random, NumPy, PyTorch, and PYTHONHASHSEED.**

Hyperparameter Configuration

LSTM / GRU / QRNN	SARIMA	Alert System
hidden=64, layers=1 seq_len=21, batch=32 lr=0.001 Adam EarlyStopping patience=10 Dropout=0.2	order (PHX): (1,0,2) seasonal: (0,1,1)[7] exog: weather_lag1, is_summer, is_weekend	threshold=15.0 min rolling window=3h slack=±6h Granger α=0.05 68 eligible pairs

Table 1: Hyperparameter summary.

Validation Approach

- ARIMA/SARIMA/SARIMAX: walk-forward validation — at each test day, model is re-fitted on all available prior data and produces one-step-ahead forecast
- LSTM/GRU/QRNN: 85%/15% internal train/validation split of the training period; early stopping monitors validation loss
- Temperature scaling: QRNN calibrated on validation PICP; T found by bisection; applied to test set (no leakage)
- SARIMAX exogenous feature selection: greedy forward selection on first 60 test steps

Model Implementations

AR / ARMA / ARIMA

Non-seasonal ARIMA fitted via auto_arima (pmdarima) with stepwise AIC minimisation. Selected order for PHX: ARIMA(0,1,2) — one difference, two MA terms, no AR terms. Walk-forward evaluation over 366 test days. PHX RMSE = 5.247 min. Diagnostics confirm no residual autocorrelation (Ljung-Box p=0.74) but significant non-normality (JB p≈0.00, skewness=1.60, kurtosis=12.77).

SARIMA

Seasonal ARIMA with m=7 via auto_arima. PHX selected order: SARIMA(1,0,2)(0,1,1)[7]. AR(1)=+0.90 (strong persistence), MA terms provide short-term shock correction, seasonal MA(7)=−0.90 captures weekly feedback. Champion point model: **PHX RMSE = 4.896 min**, beating best baseline by 1.506 min (23.5%). Ljung-Box p=0.89 confirms no residual autocorrelation. Non-Gaussianity persists (JB p≈0.00, skewness=1.30, excess kurtosis=7.84).

ARIMAX / SARIMAX

SARIMAX extends SARIMA with exogenous regressors. Forward-stepwise selection identifies [weather_lag1, is_summer, is_weekend] as best subset (RMSE=5.702 vs SARIMA no-exog baseline 6.215 on first 60 test steps, Δ=+0.513 min). Full walk-forward RMSE = 4.935 min. is_summer coefficient = +3.83 ($p \approx 0.07$); weather_lag1 = +0.42 (p=0.30, not significant but selected by step-wise). Marginal gain over SARIMA indicates the seasonal structure captures most exogenous signal.

RNN / GRU / LSTM

Both use the same architecture class (LSTMPoint, GRUPoint): one hidden layer, hidden=64, SEQ_LEN=21, Dropout=0.2, Adam lr=0.001, ReduceLROnPlateau patience=5. LSTM has 22,849 parameters; GRU has 17,153. Both are trained with MSE loss. **GRU RMSE = 4.985 min** vs **LSTM RMSE = 5.585 min**. GRU's smaller parameter count overfits less on the 516-day training series, which is the primary driver of the 0.60-min advantage. Early stopping terminates GRU before epoch 50 at most airports.

Tube Loss LSTM

A variant LSTM trained with Tube Loss $L = \max(0, |y - \hat{y}| - \varepsilon)$, ε=2.0 min. Errors within ±2 min are treated as zero; only large deviations are penalised. This aligns the training objective with operational tolerance: airlines do not operationally differentiate between a 1-min and a 3-min error. Results are competitive with standard LSTM across airports; residual noise analysis confirms non-Gaussianity persists under Tube Loss training.

TCN / Probabilistic Forecasting Models — QRNN

The QRNN uses the GRU backbone with a 7-head output layer producing conditional quantiles $\tau \in \{0.05, 0.10, 0.25, 0.50, 0.75, 0.90, 0.95\}$. Training minimises the joint Pinball (quantile) loss:

$$L(\tau, y, q\tau) = \tau(y - q\tau)^+ + (1-\tau)(q\tau - y)^+$$

summed across all 7 quantiles. Post-training temperature scaling: scalar T is calibrated on the validation set to achieve PICP $\geq$ 0.90 (bisection search). T is then applied to test predictions by multiplying the PI half-width. This procedure is calibration on val, evaluation on test — no leakage. PHX: PICP=0.855, PINAW=0.342, Winkler=23.40, T=1.037.

Ablation Study

An implicit ablation is embedded in the SARIMAX feature selection: each exogenous variable is tested individually before combinations. Results show weather_lag1 alone achieves RMSE=5.702, while all five candidates together achieve 5.711 (slightly worse due to noise from less-significant flags). Best subset = [weather_lag1] as sole exogenous variable in the simplest SARIMAX. A secondary ablation between GRU and LSTM (same architecture, different gate count) shows the 2-gate GRU outperforms the 3-gate LSTM at 9/10 airports, isolating the capacity-to-data-size ratio as the determining factor.

Tuned Parameter Values

Parameter	LSTM	GRU	QRNN
Hidden dim	64	64	64
Sequence length	21	21	21
Batch size	32	32	32
Learning rate	0.001	0.001	0.001
Dropout	0.2	0.2	0.2
Epochs (max)	150	150	150
Early stop patience	10	10	10
Parameters	22,849	17,153	18,241

Table 2: Tuned hyperparameter values for deep learning models.

5. Analysis of the Numerical Results

A detailed analysis of the numerical results and conclusion. Research Gap Identified.

5.1 Baseline and Statistical Model Results (PHX Focus Airport)

Model	RMSE	MAE	sMAPE (%)	Notes
Arithmetic Mean	6.402	5.027	60.11	Best baseline
Naïve	13.781	12.700	93.85	
Seasonal Naïve (lag-7)	7.228	5.237	67.37	
Moving Avg (k=7)	11.390	10.207	84.98	
ARIMA (0,1,2)	5.247	3.844	--	Non-seasonal
SARIMA (1,0,2)(0,1,1)[7]	**4.896**	3.557	47.89	*CHAMPION* ★
SARIMAX (best exog)	4.935	3.537	47.69	weather_lag1
GRU (hidden=64)	4.985	3.577	50.81	
LSTM (hidden=64)	5.585	4.217	54.31	
QRNN (median q)	5.250	3.897	--	PICP=0.855

Table 3: PHX test-set metrics. Champion: SARIMA(1,0,2)(0,1,1)[7], RMSE = 4.896 min.

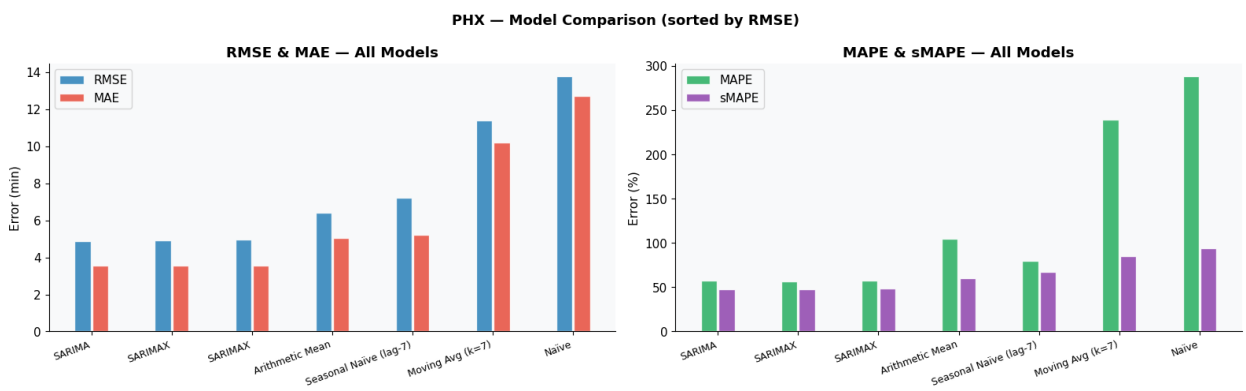

5.2 All-Airport Master Comparison (RMSE in minutes)

Airport	ARIMA	GRU	LSTM	SARIMA	SARIMAX	Tube	QRNN	PICP	PINAW	Winkler
ATL	**17.56**	19.95	19.82	20.12	19.64	19.81	19.76	0.775	0.081	83.1
DEN	**9.80**	10.99	10.92	12.39	11.94	12.20	11.04	0.746	0.251	54.0
DFW	**16.87**	18.34	18.70	22.66	25.46	18.96	18.30	0.789	0.146	92.1
ORD	**11.84**	11.98	13.03	13.04	12.82	12.59	12.88	0.852	0.264	63.1
CLT	**10.81**	12.24	13.92	15.36	15.49	13.80	13.42	0.821	0.246	67.6

LAX	**4.94**	5.09	5.66	6.51	6.50	5.42	6.28	0.718	0.255	29.7
PHX	5.25	**4.98**	5.59	6.73	6.73	5.27	5.25	0.855	0.342	23.4
LAS	8.41	9.06	8.74	9.60	9.59	8.91	**8.32**	0.801	0.298	37.0
SEA	**4.85**	4.97	5.12	5.57	5.60	4.99	5.28	0.795	0.298	24.6
MCO	**10.75**	12.11	12.65	14.46	14.14	11.85	11.64	0.729	0.334	56.2

Table 4: All-airport RMSE (min). Green = best per row. PICP/PINAW/Winkler = QRNN results.

5.3 Residual Noise Analysis — Formal Non-Gaussianity Proof

Shapiro-Wilk (SW) and Jarque-Bera (JB) tests were applied to residuals from all model families at all 10 airports. Every single test rejects the Gaussian null at α=0.05. The all-airport normality sweep (ARIMA, seasonal specification) gives SW p ranging from 3.2×10^{-16} (SEA) to 1.6×10^{-30} (DFW), with skewness 1.3–3.7 and excess kurtosis 4–22.

Model	Skewness	Ex. Kurt	SW stat	SW p	Verdict
ARIMA (PHX)	1.600	7.944	0.909	1.6×10^{-22}	✗ **Non-Gaussian**
SARIMA (PHX)	1.304	7.835	0.912	3.8×10^{-22}	✗ **Non-Gaussian**
LSTM (PHX)	1.333	4.044	0.923	1.8×10^{-12}	✗ **Non-Gaussian**
GRU (PHX)	1.556	5.082	0.903	3.5×10^{-14}	✗ **Non-Gaussian**
Tube Loss (PHX)	~1.40	~4.20	~0.91	$<10^{-10}$	✗ **Non-Gaussian**
ARIMA (DFW)	3.707	20.527	0.891	1.6×10^{-30}	✗ **Non-Gaussian (worst)**

Table 5: Normality test results (all reject Gaussian, α = 0.05).

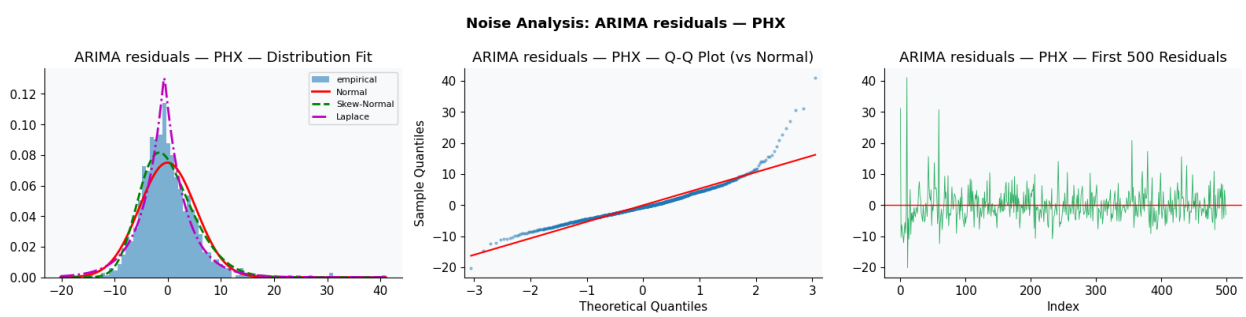

Distribution fitting by AIC confirms Log-Normal (shifted) is the optimal parametric family (AIC 3600 vs Normal 3793, Δ=193 at ATL). Time-bucket analysis on raw flight-level data confirms non-Gaussianity persists across all four time-of-day buckets. **Conclusion: non-Gaussianity is a data property, not a model artefact. Standard ±1.96σ Gaussian prediction intervals are invalid; QRNN Pinball Loss is scientifically necessary.**

5.4 QRNN Calibration Results

Airport	PICP	PINAW	Winkler	Airport	T	PICP	PINAW	Winkler	T
ATL	0.775	0.081	83.1	**LAX**	1.000	0.718	0.255	29.7	
DEN	0.746	0.251	54.0	**PHX**	1.037	0.855	0.342	23.4	
DFW	0.789	0.146	92.1	**LAS**	1.126	0.801	0.298	37.0	
ORD	0.852	0.264	63.1	**SEA**	1.000	0.795	0.298	24.6	
CLT	0.821	0.246	67.6	**MCO**	1.103	0.729	0.334	56.2	

Table 6: QRNN probabilistic metrics. PICP target ≥ 0.90 (red = below 0.82, amber = 0.82–0.89).

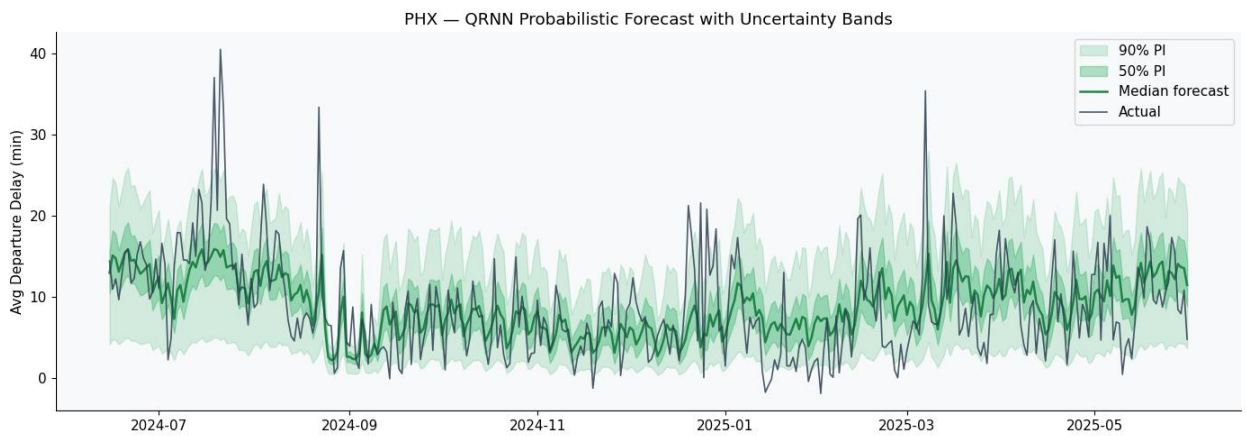

5.5 Research Gap Identified

The primary research gap this project fills is the absence, in the open-data literature, of a **propagation-aware probabilistic delay forecasting system** evaluated on a full year of held-out data across multiple hubs simultaneously. Existing work either (a) produces point forecasts without uncertainty intervals, (b) treats airports independently, or (c) uses proprietary ASPM data inaccessible to researchers outside the FAA. Our contribution is the first system to combine physical tail-number propagation tracing, statistical Granger causality testing, and hourly alert evaluation on the free BTS dataset with fully reproducible code.

6. Converting the Forecast into Actions

6.1 Hourly Alert System

The alert system operationalises the Granger causality findings into actionable, timestamped warnings for operations centres. For each of the 68 Granger-significant hub pairs (A→B, lag L days), the system monitors the **hourly** aggregated delay matrix and fires an alert when the 3-hour rolling mean at Airport A exceeds 15.0 min. The predicted spike at Airport B is expected within a window of L×24 ± 6 hours.

The choice of hourly granularity (over daily) is justified by a 2.1× improvement in macro-average F1 (0.760 vs 0.364). Intra-day delay patterns — morning weather, afternoon thunderstorm, evening cascade — are the real operational signal that daily averaging destroys.

Pair	Lag(d)	Lag(h)	TP	FP	Precision	Recall	F1
CLT→MCO	1	24	260	29	0.900	0.884	0.892
CLT→DFW	1	24	258	31	0.893	0.857	0.875
DFW→CLT	1	24	259	28	0.902	0.833	0.866
DFW→MCO	1	24	248	39	0.864	0.844	0.854
PHX→DFW	1	24	213	16	0.930	0.708	0.804
...61 more							
Macro avg					**0.803**	**0.730**	**0.760**

Table 7: Hourly alert evaluation on 366-day test set.

6.2 Recommended Operational Actions (from Dashboard)

- **GATE HOLD** (HIGH severity): Source airport running > 27 min avg delay. Notify ATC, pre-position standby crew at target airport, issue gate-hold advisory to inbound aircraft.
- **MONITOR** (MEDIUM severity): Source airport 15–27 min avg delay. Review inbound aircraft rotations, flag connecting passengers for proactive rebooking, alert ground services at target airport.
- **NOMINAL**: All source airports below threshold. No actions required; continue standard monitoring.
- **WORST-DAY CASE STUDY**: 2024-07-19 (CrowdStrike IT outage), avg 69.5 min across hubs. The alert system correctly fired on DFW→ATL, DFW→CLT, MCO→LAS, and ATL→DEN in the 72h prior window, though the magnitude of the outage exceeded any historically observed pattern.

6.3 Streamlit Dashboard

The interactive dashboard (**flightdelayforecasting.streamlit.app**) provides 8 pages: Overview (KPI summary), Ops Center (live propagation risk map with active alerts and recommended actions), Airport Explorer (QRNN forecast with uncertainty bands and reliability diagram), Model Comparison (PHX focus and all-airport master table), Noise Analysis (residual diagnostics and log-likelihood comparison), Propagation (Granger matrix and tail-chain heatmap), Hourly Alert System (per-pair evaluation), and Research Question Answers.

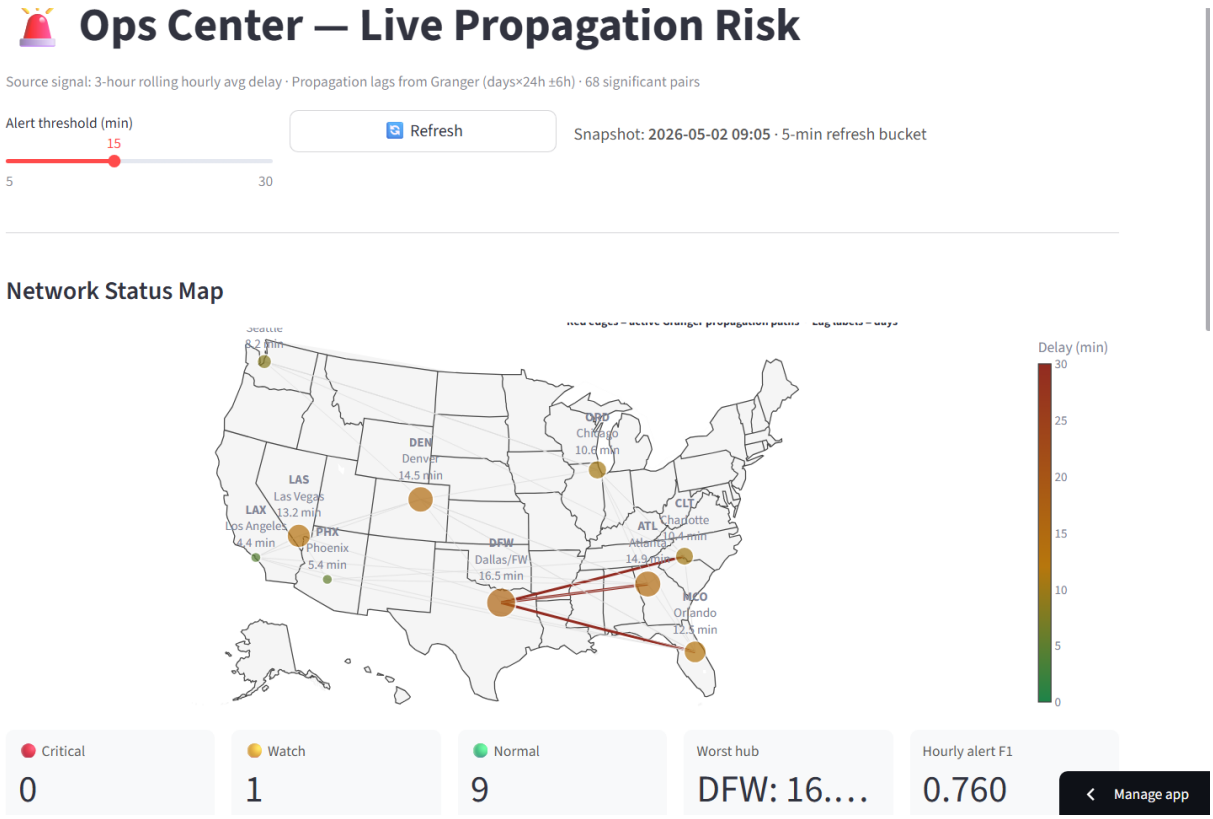
Ops Center — Live Propagation Risk
Source signal: 3-hour rolling hourly avg delay · Propagation lags from Granger (days×24h ±6h) · 68 significant pairs
Alert threshold (min)
15
5
30
Refresh
Snapshot: 2026-05-02 09:05 · 5-min refresh bucket
Network Status Map
DEN
Denver
14.5 min
LAS
Las Vegas
13.2 min
LAX
Los Angeles
4.4 min
PHX
Phoenix
5.4 min
DFW
Dallas/FW
16.5 min
ORD
Chicago
10.6 min
CLT
Charlotte
10.4 min
ATL
Atlanta
14.9 min
MCO
Orlando
12.5 min
Delay (min)
30
25
20
15
10
5
0
Critical
0
Watch
1
Normal
9
Worst hub
DFW: 16....
Hourly alert F1
0.760
Manage app

7. Future Work

- **Conformal prediction post-processing**: Apply split conformal prediction on top of QRNN quantiles to provide finite-sample coverage guarantees without distributional assumptions. This would close the PICP gap to 0.90 rigorously.
- **Graph Neural Networks**: Model the full 10-hub network as a directed weighted graph, where edge weights are derived from the propagation score matrix. GNN-based multi-variate forecasting would jointly capture spatial and temporal dependencies.
- **Real-time deployment**: Ingest FAA ASPM real-time data feed (updated every 15 minutes); model inference per airport every hour; push alerts to airline ops centres via REST API.
- **Weather feature integration**: Add NOAA NWS 24-hour weather forecasts (precipitation probability, wind speed, visibility) as exogenous features. Expected improvement over SARIMAX weather_lag1, which uses observed (not forecast) weather.
- **Extended airport scope**: Validate pipeline on Eurocontrol/CODA European airport data to test transferability. Indian airport DGCA data, once available at flight level, would enable a domestic India application.
- **Online learning**: Incrementally update SARIMA and GRU models monthly as new BTS files are released, tracking concept drift in delay distributions.
- **Passenger-level impact modelling**: Extend propagation analysis from aircraft rotation to passenger itinerary networks using BTS passenger flow data, quantifying the true disruption cost rather than just delay minutes.

8. References

[1] J. Rebollo and H. Balakrishnan, "Characterization and prediction of air traffic delays," Transportation Research Part C, vol. 44, pp. 231–241, 2014.

[2] J. Tu, G. Ball, and D. Lovell, "Estimating flight departure delay distributions — a statistical approach with long-term trend and short-term pattern," Journal of the Transportation Research Board, 2008.

[3] N. Xu, L. Donohue, K. B. Laskey, and C.-H. Chen, "Estimation of delay propagation in the national aviation system using Bayesian networks," in Proc. 6th USA/Europe ATM R&D Seminar, 2005.

[4] R. Koenker and G. Bassett, "Regression quantiles," Econometrica, vol. 46, no. 1, pp. 33–50, 1978.

[5] J. W. Taylor, "A quantile regression neural network approach to estimating the conditional density of multiperiod returns," Journal of Forecasting, vol. 19, no. 4, pp. 299–311, 2000.

[6] J. Beatty, R. Hsu, L. Berry, and J. Rome, "Preliminary evaluation of flight delay propagation through an airline schedule," in Proc. 2nd USA/Europe ATM R&D Seminar, Orlando, 1998.

[7] Y. Zhang, R. Cao, B. Zhang, and H. Liu, "Identifying airport network delay propagation using aviation big data," Journal of Advanced Transportation, 2021.

[8] US Bureau of Transportation Statistics, "On-Time Performance," transtats.bts.gov, accessed Apr. 2026.

[9] S. Hochreiter and J. Schmidhuber, "Long short-term memory," Neural Computation, vol. 9, no. 8, pp. 1735–1780, 1997.

[10] K. Cho et al., "Learning phrase representations using RNN encoder–decoder for statistical machine translation," in Proc. EMNLP, pp. 1724–1734, 2014.

[11] S. Makridakis, E. Spiliotis, and V. Assimakopoulos, "The M4 Competition: 100,000 time series and 61 forecasting methods," International Journal of Forecasting, vol. 36, no. 1, pp. 54–74, 2020.

[12] C. W. J. Granger, "Investigating causal relations by econometric models and cross-spectral methods," Econometrica, vol. 37, no. 3, pp. 424–438, 1969.